

Distributed Constraint Optimization

Gauthier Picard

MINES Saint-Étienne
LaHC UMR CNRS 5516

gauthier.picard@emse.fr

— Some contents taken from OPTMAS 2011 and OPTMAS-DCR 2014 Tutorials—

Contents

Introduction

- Constraint Optimization Problems
- DCOP Framework
- Application Domains

Complete Algorithms for DCOP

- Asynchronous Distributed Optimisation (ADOPT)
- Dynamic Programming Optimization Protocol (DPOP)

Approximate Algorithms for DCOP

- Distributed Stochastic Search Algorithm (DSA)
- Maximum Gain Message (MGM-1)

Synthesis

- Panorama

Constraint Optimization Problems

Sometimes satisfaction is not possible

- Overconstrained problem
- Solution is not binary

Switch from satisfaction to optimization

- Minimizing the number of violated constraints
- Minimizing the cost of violated constraints
- Maximizing the overall utility of the system
- ...

DCOP Framework

Motivations

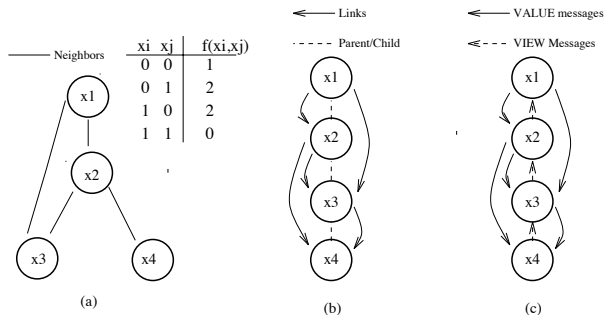
- In dynamic and complex environments not all constraints can be satisfied completely
- Satisfaction → **Optimisation** (combinatorial)
 - ▶ ex: minimizing the number of unchecked constraints, minimizing the sum of the costs of violated constraints, etc.

Definition (DCOP)

A DCOP is a DCSP $\langle A, X, D, C, \phi \rangle$ with

- a cost function $f_{ij} : D_i \times D_j \mapsto \mathbb{N} \cup \infty$ for each pair x_i, x_j
- an objective function $F : D \mapsto \mathbb{N} \cup \infty$ evaluating an assignment \mathcal{A} with $f_{ij}(d_i, d_j)$ for each pair x_i, x_j

DCOP Framework (cont.)



Objective Function

$$F(\mathcal{A}) = \sum_{x_i, x_j \in X} f_{ij}(d_i, d_j) \text{ where } x_i \leftarrow d_i \text{ and } x_i \leftarrow d_i \text{ in } \mathcal{A}$$

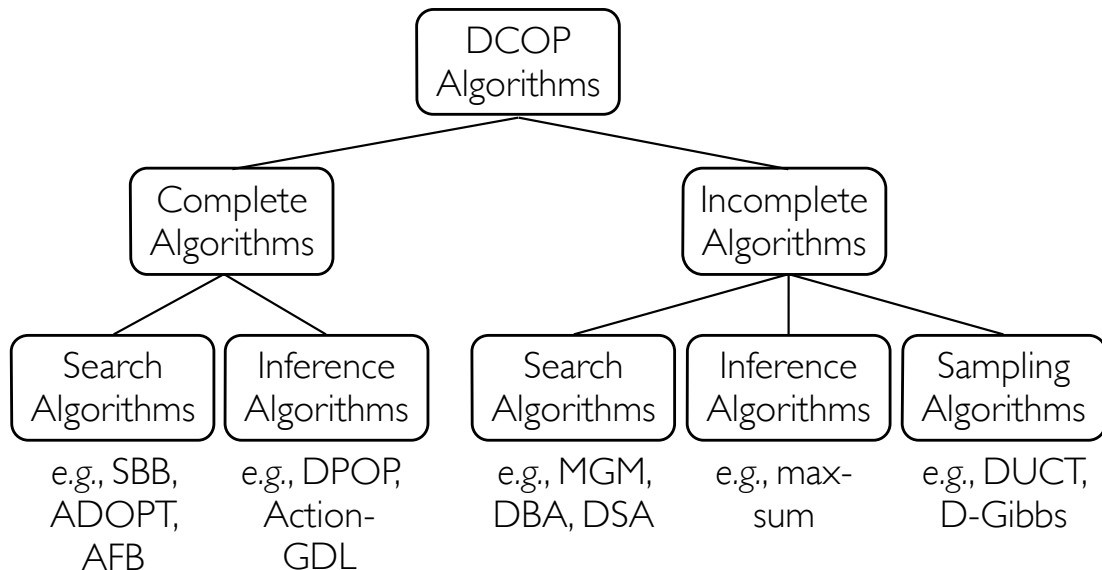
In figure (a):

- $F(\{(x_1, 0), (x_2, 0), (x_3, 0), (x_4, 0)\}) = 4$
- $F(\{(x_1, 1), (x_2, 1), (x_3, 1), (x_4, 1)\}) = 0$



and

$$\mathcal{A}^* = \{(x_1, 1), (x_2, 1), (x_3, 1), (x_4, 1)\}$$

DCOP Algorithms



Application Domains



Contents

Introduction

Complete Algorithms for DCOP

- Asynchronous Distributed Optimisation (ADOPT)

- Dynamic Programming Optimization Protocol (DPOP)

Approximate Algorithms for DCOP

Synthesis


Asynchronous Distributed Optimisation (ADOPT) [MODI et al., 2005]

ADOPT: DFS tree (pseudotree)

ADOPT assumes that agents are arranged in a DFS tree:

- constraint graph \rightarrow rooted graph (select a node as root)
- some links form a tree / others are backedges
- two constrained nodes must be in the same path to the root by tree links (same branch)

Every graph admits a DFS tree: DFS graph traversal



ADOPT Features

- Asynchronous algorithm
- Each time an agent receives a message:
 - ▶ Processes it (the agent may take a new value)
 - ▶ Sends VALUE messages to its children and pseudochildren
 - ▶ Sends a COST message to its parent
- Context: set of (variable value) pairs (as ABT agent view) of ancestor agents (in the same branch)
- Current context:
 - ▶ Updated by each VALUE message
 - ▶ If current context is not compatible with some child context, the later is initialized (also the child bounds)

ADOPT Procedures

Initialize

```
(1) threshold  $\leftarrow 0$ ; CurrentContext  $\leftarrow \{\}$ ;
(2) forall  $d \in D_i, x_i \in \text{Children}$  do
(3)   lb( $d, x_i$ )  $\leftarrow 0$ ; t( $d, x_i$ )  $\leftarrow 0$ ;
(4)   ub( $d, x_i$ )  $\leftarrow \text{Inf}$ ; context( $d, x_i$ )  $\leftarrow \{\}$ ; enddo;
(5)  $d_i \leftarrow d$  that minimizes LB( $d$ );
(6) backTrack;
```

```
when received (THRESHOLD,  $t$ , context)
(7) if context compatible with CurrentContext:
(8)   threshold  $\leftarrow t$ ;
(9)   maintainThresholdInvariant;
(10)  backTrack; endif;
```

```
when received (TERMINATE, context)
(11) record TERMINATE received from parent;
(12) CurrentContext  $\leftarrow \text{context}$ ;
(13) backTrack;
```

```
when received (VALUE, ( $x_j, d_j$ ))
(14) if TERMINATE not received from parent:
(15)   add ( $x_j, d_j$ ) to CurrentContext;
(16)   forall  $d \in D_i, x_i \in \text{Children}$  do
(17)     if context( $d, x_i$ ) incompatible with CurrentContext:
(18)       lb( $d, x_i$ )  $\leftarrow 0$ ; t( $d, x_i$ )  $\leftarrow 0$ ;
(19)       ub( $d, x_i$ )  $\leftarrow \text{Inf}$ ; context( $d, x_i$ )  $\leftarrow \{\}$ ; endif; enddo;
(20)   maintainThresholdInvariant;
(21)   backTrack; endif;
```

```
when received (COST,  $x_k$ , context, lb, ub)
(22)  $d \leftarrow$  value of  $x_i$  in context;
(23) remove ( $x_i, d$ ) from context;
(24) if TERMINATE not received from parent:
(25)   forall ( $x_j, d_j$ )  $\in$  context and  $x_j$  is not my neighbor do
(26)     add ( $x_j, d_j$ ) to CurrentContext; enddo;
(27)   forall  $d' \in D_i, x_i \in \text{Children}$  do
(28)     if context( $d', x_i$ ) incompatible with CurrentContext:
(29)       lb( $d', x_i$ )  $\leftarrow 0$ ; t( $d', x_i$ )  $\leftarrow 0$ ;
(30)       ub( $d', x_i$ )  $\leftarrow \text{Inf}$ ; context( $d', x_i$ )  $\leftarrow \{\}$ ; endif; enddo; endif;
(31) if context compatible with CurrentContext:
(32)   lb( $d, x_k$ )  $\leftarrow \text{lb}$ ;
(33)   ub( $d, x_k$ )  $\leftarrow \text{ub}$ ;
(34)   context( $d, x_k$ )  $\leftarrow \text{context}$ ;
(35)   maintainChildThresholdInvariant;
(36)   maintainThresholdInvariant; endif;
(37) backTrack;
```

procedure backTrack

```
(38) if threshold == UB:
(39)    $d_i \leftarrow d$  that minimizes UB( $d$ );
(40) else if LB( $d_i$ ) > threshold:
(41)    $d_i \leftarrow d$  that minimizes LB( $d$ ); endif;
(42) SEND (VALUE, ( $x_i, d_i$ ))
(43)   to each lower priority neighbor;
(44) maintainAllocationInvariant;
(45) if threshold == UB:
(46)   if TERMINATE received from parent
(47)     or  $x_i$  is root:
(48)     SEND (TERMINATE,
(49)       CurrentContext  $\cup \{(x_i, d_i)\}$ )
(50)       to each child;
(51)     Terminate execution; endif; endif;
(52) SEND (COST,  $x_i$ , CurrentContext, LB, UB)
    to parent;
```

Algorithm 1: ADOPT Procedures

ADOPT Messages

- **value**($parent \rightarrow children \cup pseudochildren, a$): parent informs descendants that it has taken value a
- **cost**($child \rightarrow parent, lowerbound, upperbound, context$): child informs parent of the best cost of its assignment; attached context to detect obsolescence
- **threshold**($parent \rightarrow child, t$): minimum cost of solution in child is at least t
- **termination**($parent \rightarrow children$): sent when $LB = UB$

ADOPT Data Structures

1. **Current context** (agent view): values of higher priority constrained agents

x_i	x_j	...
a	c	...


2. **Bounds** (for each value, child)

- ▶ lower bounds
- ▶ upper bounds
- ▶ thresholds
- ▶ contexts

x_j	a	b	c	d
$lb(x_k)$	3	0	0	0
$ub(x_k)$	∞	∞	∞	∞
$th(x_k)$	1	0	0	0
$context(x_k)$				

- Stored contextes must be active: $context \in currentcontext$
- If a context becomes no active, it is removed ($lb \leftarrow 0, th \leftarrow 0, ub \leftarrow \infty$)


ADOPT Bounds



ADOPT Bounds

$\delta(value) = \text{cost with higher agents}$


$$\delta(b) = \sum_{i \in \text{curctx}} c_{ij}(a, b)$$



ADOPT Bounds

$\delta(\text{value}) = \text{cost with higher agents}$

$$\delta(b) = \sum_{i \in \text{curctx}} c_{ij}(a, b)$$




$$OPT(x_j, ctx) = \min_{d \in d_j} \delta(d) + \sum_{x_k \in \text{children}} OPT(x_k, ctx \cup (x_j, d))$$

ADOPT Bounds

$\delta(\text{value}) = \text{cost with higher agents}$

$$\delta(b) = \sum_{i \in \text{curctx}} c_{ij}(a, b)$$




$$OPT(x_j, ctx) = \min_{d \in d_j} \delta(d) + \sum_{x_k \in \text{children}} OPT(x_k, ctx \cup (x_j, d))$$

ADOPT Bounds

$\delta(value) = \text{cost with higher agents}$

$$\delta(b) = \sum_{i \in curctx} c_{ij}(a, b)$$




$[lb_k, ub_k] = \text{cost of lower agents}$

$$OPT(x_j, ctx) = \min_{d \in d_j} \delta(d) + \sum_{x_k \in children} OPT(x_k, ctx \cup (x_j, d))$$

ADOPT Bounds

$\delta(\text{value}) = \text{cost with higher agents}$

$$\delta(b) = \sum_{i \in \text{curctx}} c_{ij}(a, b)$$



$[lb_k, ub_k] = \text{cost of lower agents}$

$$OPT(x_j, ctx) = \min_{d \in d_j} \delta(d) +$$

$$\sum_{x_k \in \text{children}} OPT(x_k, ctx \cup (x_j, d))$$

$$LB(b) = \delta(b) + \sum_{x_k \in \text{children}} lb(b, x_k)$$

$$LB = \min_{b \in d_j} LB(b)$$

$$UB(b) = \delta(b) + \sum_{x_k \in \text{children}} ub(b, x_k)$$

$$UB = \min_{b \in d_j} UB(b)$$

ADOPT Value Assignment

- An ADOPT agent takes the value with minimum LB
- Eager behavior:
 - ▶ Agents may constantly change value
 - ▶ Generates many context changes
- Threshold:
 - ▶ lower bound of the cost that children have from previous search
 - ▶ parent distributes threshold among children
 - ▶ incorrect distribution does not cause problems: the child with minor allocation would send a COST to the parent later, and the parent will rebalance the threshold distribution

ADOPT Properties

- For any x_i , $LB \leq OPT(x_l, ctx) \leq UB$
- For any x_i , its threshold reaches UB
- For any x_i , its final threshold is equal to $OPT(x_l, ctx)$

→ **ADOPT terminates with the optimal solution**


ADOPT Example

- 4 variables (4 agents) x_1, x_2, x_3 and x_4 with $D = \{a, b\}$


- 4 binary identical cost functions

x_i	x_j	cost
a	a	1
a	b	2
b	a	2
b	b	0



- Constraint graph






ADOPT Example (cont.)




ADOPT Example (cont.)




ADOPT Example (cont.)




ADOPT Example (cont.)



ADOPT Example (cont.)



ADOPT Example (cont.)



Dynamic Programming Optimization Protocol (DPOP) [PETCU and FALTINGS, 2005]

3-phase distributed algorithm

PHASES	MESSAGES
1. DFS Tree construction	token passing
2. Utility phase: from leaves to root	util (child \rightarrow parent, constraint table [-child])
3. Value phase: from root to leaves	value (parent \rightarrow children \cup pseudochildren, parent value)

DFS Tree Phase


■ Distributed DFS graph traversal: token, ID, $neighbors(X)$

1. X owns the token: adds its own ID and sends it in turn to each of its neighbors, which become children
2. Y receives the token from X : it marks X as visited. First time Y receives the token then $parent(Y) = X$. Other IDs in token which are also $neighbors(Y)$ are **pseudoparent**. If Y receives token from neighbor W to which it was never sent, W is pseudochild.
3. When all $neighbors(X)$ visited, X removes its ID from token and sends it to $parent(X)$.


■ A node is selected as root, which starts

■ When all neighbors of root are visited, the DFS traversal ends


DFS Tree Phase: Example



DFS Tree Phase: Example



x_1 parent of x_2




x_2 parent of x_3


x_1 pseudoparent of x_3

DFS Tree Phase: Example


root



x_1 parent of x_2




x_2 parent of x_3
 x_1 pseudoparent of x_3




x_3 parent of x_4
 x_3 pseudoparent of x_1


DFS Tree Phase: Example




x_1 parent of x_2



x_2 parent of x_3
 x_1 pseudoparent of x_3



x_3 parent of x_4
 x_3 pseudoparent of x_1



Util Phase

Agent X :


- receives from each child Y_i a cost function: $C(Y_i)$
- combines (adds, joins) all these cost functions with the cost functions with $parent(X)$ and $pseudoparents(X)$
- projects X out of the resulting cost function, and sends it to $parent(X)$

From the leaves to the root


Util Phase: Example



Util Phase: Example




Util Phase: Example



All value combinations
Costs are the sum of applicable costs

Util Phase: Example




Value Phase


1. The root finds the **value that minimizes the received cost function** in the util phase, and informs its descendants (children \cup pseudochildren)
2. Each agent **waits to receive** the value of its parent / pseudoparents
3. Keeping fixed the value of parent/pseudoparents, finds **the value that minimizes the received cost function** in the Util phase
4. Informs of this value to its children/pseudochildren

This process starts at the root and ends at the leaves


DTREE : DPOP for DCOPs without backedges




DTREE : DPOP for DCOPs without backedges




DTREE : DPOP for DCOPs without backedges




DTREE : DPOP for DCOPs without backedges




DTREE : DPOP for DCOPs without backedges




DTREE : DPOP for DCOPs without backedges




DTREE : DPOP for DCOPs without backedges



DTREE : DPOP for DCOPs without backedges




DTREE : DPOP for DCOPs without backedges




Optimal solution:

- linear number of messages
- message size: linear


DPOP for any DCOP



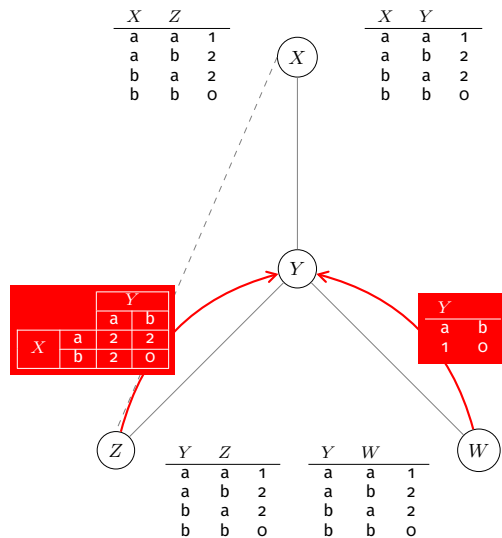
DPOP for any DCOP



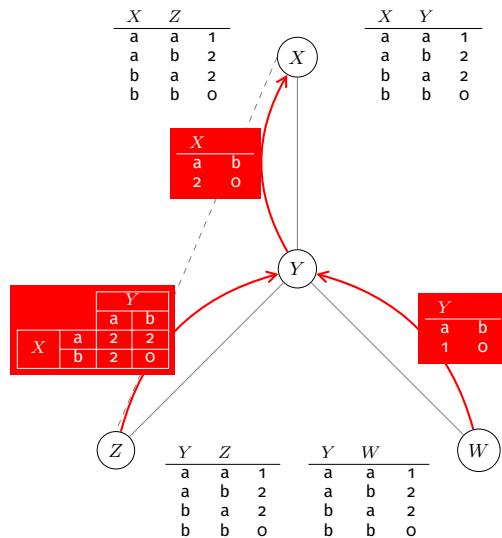
DPOP for any DCOP




DPOP for any DCOP




DPOP for any DCOP




DPOP for any DCOP




DPOP for any DCOP



DPOP for any DCOP



DPOP for any DCOP



Optimal solution:

- linear number of messages
- message size: exponential

Contents

Introduction

Complete Algorithms for DCOP

Approximate Algorithms for DCOP

Distributed Stochastic Search Algorithm (DSA)

Maximum Gain Message (MGM-1)

Synthesis

Approximate Algorithms for DCOPs

Complete algorithms

- e.g. ADOPT [MODI et al., 2005] and DPOP [PETCU and FALTINGS, 2005]
 - ✓ complete
 - ✗ slow

Approximate algorithms exist (fast, but sub-optimal in many case)

- Search algorithms
 - ▶ DBA [YOKOO, 2001], DSA [ZHANG et al., 2005], MGM [MAHESWARAN et al., 2004]
- Inference algorithms
 - ▶ Max-sum [FARINELLI et al., 2008]

Why Approximate Algorithms

■ Motivations

- ▶ Often optimality in practical applications is not achievable
- ▶ Fast good enough solutions are all we can have

■ Example – Graph coloring

- ▶ Medium size problem (about 20 nodes, three colors per node)
- ▶ Number of states to visit for optimal solution in the worst case $3^{20} = 3M$ states

■ Key problem

- ▶ Provides guarantees on solution quality

Exemplar Application: Surveillance

■ Event Detection

- ▶ Vehicles passing on a road

■ Energy Constraints


- ▶ Sense/Sleep modes
- ▶ Recharge when sleeping

■ Coordination

- ▶ Activity can be detected by single sensor
- ▶ Roads have different traffic loads

■ Aim


- ▶ Focus on road with more traffic load



Centralized Local Greedy approaches

■ Greedy local search


- ▶ Start from random solution
- ▶ Do local changes if global solution improves
- ▶ Local: change the value of a subset of variables, usually one



Centralized Local Greedy approaches

■ Greedy local search


- ▶ Start from random solution
- ▶ Do local changes if global solution improves
- ▶ Local: change the value of a subset of variables, usually one



Centralized Local Greedy approaches

■ Greedy local search


- ▶ Start from random solution
- ▶ Do local changes if global solution improves
- ▶ Local: change the value of a subset of variables, usually one



Centralized Local Greedy approaches

■ Greedy local search


- ▶ Start from random solution
- ▶ Do local changes if global solution improves
- ▶ Local: change the value of a subset of variables, usually one



Centralized Local Greedy approaches

■ Greedy local search


- ▶ Start from random solution
- ▶ Do local changes if global solution improves
- ▶ Local: change the value of a subset of variables, usually one



Centralized Local Greedy approaches

■ Problems


- ▶ Local minima
- ▶ Standard solutions: Random Walk, Simulated Annealing



Centralized Local Greedy approaches

■ Problems


- ▶ Local minima
- ▶ Standard solutions: Random Walk, Simulated Annealing



Centralized Local Greedy approaches

■ Problems


- ▶ Local minima
- ▶ Standard solutions: Random Walk, Simulated Annealing



Centralized Local Greedy approaches


■ Problems

- ▶ Local minima
- ▶ Standard solutions: Random Walk, Simulated Annealing




Distributed Local Greedy approaches

- Local knowledge
- Parallel execution
 - ▶ A greedy local move might be harmful/useless
 - ▶ Need coordination




Distributed Local Greedy approaches

- Local knowledge
- Parallel execution
 - ▶ A greedy local move might be harmful/useless
 - ▶ Need coordination




Distributed Local Greedy approaches

- Local knowledge
- Parallel execution
 - ▶ A greedy local move might be harmful/useless
 - ▶ Need coordination



Distributed Local Greedy approaches


- Local knowledge
- Parallel execution
 - ▶ A greedy local move might be harmful/useless
 - ▶ Need coordination




Distributed Stochastic Search Algorithm (DSA) [ZHANG et al., 2005]

- Greedy local search with activation probability to mitigate issues with parallel executions
- DSA-1: change value of one variable at time
- Initialize agents with a random assignment and communicate values to neighbors
- Each agent:
 - ▶ Generates a random number and execute only if rnd less than activation probability
 - ▶ When executing changes value maximizing local gain
 - ▶ Communicate possible variable change to neighbors


DSA-1: Execution Example




DSA-1: Execution Example




DSA-1: Execution Example



DSA-1: Execution Example



DSA-1: Execution Example




DSA-1: Discussion

- Extremely “cheap” (computation/communication)
- Good performance in various domains
 - ▶ e.g. target tracking [FITZPATRICK and MEERTENS, 2003; ZHANG et al., 2003]
 - ▶ Shows an anytime property (not guaranteed)
 - ▶ Benchmarking technique for coordination
- Problems
 - ▶ Activation probability must be tuned [ZHANG et al., 2003]
 - ▶ No general rule, hard to characterise results across domains


Maximum Gain Message (MGM-1) [MAHESWARAN et al., 2004]

- Coordinate to decide who is going to move
 - ▶ Compute and exchange possible gains
 - ▶ Agent with maximum (positive) gain executes
- Analysis
 - ▶ Empirically, similar to DSA
 - ▶ More communication (but still linear)
 - ▶ **No Threshold to set**
 - ▶ **Guaranteed to be monotonic** (Anytime behavior)


MGM-1: Example




MGM-1: Example




MGM-1: Example




MGM-1: Example




MGM-1: Example




MGM-1: Example




MGM-1: Example




MGM-1: Example




MGM-1: Example




MGM-1: Example



MGM-1: Example



MGM-1: Example



To sum up on local greedy approaches

- Exchange local values for variables
 - ▶ Similar to search based methods (e.g. ADOPT)
- Consider only local information when maximizing
 - ▶ Values of neighbors
- Anytime behaviors
- Could result in very bad solutions


GDL-based approaches

- Generalized Distributive Law [AJI and McELIECE, 2000]
 - ▶ Unifying framework for inference in Graphical models
 - ▶ Builds on basic mathematical properties of semi-rings
 - ▶ Widely used in Info theory, Statistical physics, Probabilistic models
- Max-sum
 - ▶ DCOP settings: maximise social welfare

	K	"(+, 0)"	"(·, 1)"	short name
1.	A	(+, 0)	(·, 1)	
2.	$A[x]$	(+, 0)	(·, 1)	
3.	$A[x, y, \dots]$	(+, 0)	(·, 1)	
4.	$[0, \infty)$	(+, 0)	(·, 1)	sum-product
5.	$(0, \infty]$	(min, ∞)	(·, 1)	min-product
6.	$[0, \infty)$	(max, 0)	(·, 1)	max-product
7.	$(-\infty, \infty]$	(min, ∞)	(+, 0)	min-sum
8.	$[-\infty, \infty)$	(max, $-\infty$)	(+, 0)	max-sum
9.	$\{0, 1\}$	(OR, 0)	(AND, 1)	Boolean
10.	2^S	(\cup , \emptyset)	(\cap , S)	
11.	Λ	(\vee , 0)	(\wedge , 1)	
12.	Λ	(\wedge , 1)	(\vee , 0).	


Max-Sum

Agents iteratively compute local functions that depend only on the variable they control




Max-Sum

Agents iteratively compute local functions that depend only on the variable they control



Max-Sum


Agents iteratively compute local functions that depend only on the variable they control



$$m_{1 \rightarrow 2}(x_2) = \max_{x_1} (F_{12}(x_1, x_2) + m_{4 \rightarrow 1}(x_1))$$

Max-Sum

Agents iteratively compute local functions that depend only on the variable they control




Shared constraint

$$m_{1 \rightarrow 2}(x_2) = \max_{x_1} (\text{Shared constraint} + m_{4 \rightarrow 1}(x_1))$$

Max-Sum

Agents iteratively compute local functions that depend only on the variable they control




Shared constraint

$$m_{1 \rightarrow 2}(x_2) = \max_{x_1} (\text{[blue box]} + \text{[blue box]})$$

All incoming
messages
except x_2

Max-Sum

Agents iteratively compute local functions that depend only on the variable they control




Shared constraint

$$m_{1 \rightarrow 2}(x_2) = \max_{x_1} (\text{[blue box]} + \text{[blue box]})$$

All incoming messages except x_2

Max-Sum

Agents iteratively compute local functions that depend only on the variable they control



Shared constraint


$$m_{1 \rightarrow 2}(x_2) = \max_{x_1} (\text{[redacted]} + \text{[redacted]})$$

$$z_1(x_1) = m_{4 \rightarrow 1}(x_1) + m_{2 \rightarrow 1}(x_1)$$

All incoming
messages
except x_2

Max-Sum

Agents iteratively compute local functions that depend only on the variable they control



Shared constraint


$$m_{1 \rightarrow 2}(x_2) = \max_{x_1} (\text{[redacted]} + \text{[redacted]})$$

$$z_1(x_1) = m_{4 \rightarrow 1}(x_1) + m_{2 \rightarrow 1}(x_1)$$

All incoming
messages
except x_2

Max-Sum

Agents iteratively compute local functions that depend only on the variable they control



Shared constraint

$$m_{1 \rightarrow 2}(x_2) = \max_{x_1} (\boxed{\phantom{m_{4 \rightarrow 1}(x_1)}} + \boxed{})$$


$$z_1(x_1) = \boxed{\phantom{m_{4 \rightarrow 1}(x_1)}}$$

All incoming

All incoming
messages
except x_2

Max-Sum

Agents iteratively compute local functions that depend only on the variable they control



Shared constraint

$$m_{1 \rightarrow 2}(x_2) = \max_{x_1} (\text{[blue box]} + \text{[blue box]})$$

$$z_1(x_1) = \text{[blue box]}$$

All incoming

All incoming
messages
except x_2


Max-Sum on acyclic graphs

■ Max-sum Optimal on acyclic graphs

- ▶ Different branches are independent
- ▶ Each agent can build a correct estimation of its contribution to the global problem (z functions)

■ Message equations very similar to Util messages in DPOP



- ▶ Sum messages from children and shared constraint
- ▶ Maximize out agent variable
- ▶ GDL generalizes DPOP [VINYALS et al., 2011]



$$m_{1 \rightarrow 2}(x_2) = \max_{x_1} (F_{12}(x_1, x_2) + m_{4 \rightarrow 1}(x_1))$$


Max-sum Performance

- **Good performance on loopy networks** [FARINELLI et al., 2008]
 - ▶ When it converges very good results
 - ▶ Interesting results when only one cycle [WEISS, 2000]
 - ▶ We could remove cycle but pay an exponential price (see DPOP)



Max-Sum for low power devices

- Low overhead
 - ▶ Msgs number/size
- Asynchronous computation
 - ▶ Agents take decisions whenever new messages arrive
- Robust to message loss



Contents

Introduction

Complete Algorithms for DCOP

Approximate Algorithms for DCOP

Synthesis

Panorama

Panorama

Algorithm	Type	Memory	Messages	Remarks
ADOPT	COP	Polynomial	Exponential	Complete
DPOP	COP	Exponential	Linear	Complete
DSA	COP	Linear	?	Not complete
MGM	COP	Linear	?	Not complete
Max-Sum	COP	Exponential	Linear on acyclic	Complete on trees

Table: DCOP algorithms

References



AJI, S.M. and R.J. McELIECE (2000). “The generalized distributive law”. In: *Information Theory, IEEE Transactions on* 46.2, pp. 325–343. ISSN: 0018-9448. DOI: 10.1109/18.825794.



FARINELLI, A., A. ROGERS, A. PETCU, and N. R. JENNINGS (2008). “Decentralised Coordination of Low-power Embedded Devices Using the Max-sum Algorithm”. In: *Proceedings of the 7th International Joint Conference on Autonomous Agents and Multiagent Systems - Volume 2. AAMAS '08*. International Foundation for Autonomous Agents and Multiagent Systems, pp. 639–646. ISBN: 978-0-9817381-1-6.



FITZPATRICK, Stephen and Lambert MEERTENS (2003). “Distributed Coordination through Anarchic Optimization”. In: *Distributed Sensor Networks: A Multiagent Perspective*. Ed. by Victor LESSER, Charles L. ORTIZ, and Milind TAMBE. Boston, MA: Springer US, pp. 257–295. ISBN: 978-1-4615-0363-7.



MAHESWARAN, R.T., J.P. PEARCE, and M. TAMBE (2004). “Distributed Algorithms for DCOP: A Graphical-Game-Based Approach”. In: *Proceedings of the 17th International Conference on Parallel and Distributed Computing Systems (PDCS), San Francisco, CA*, pp. 432–439.



MODI, P. J., W. SHEN, M. TAMBE, and M. YOKOO (2005). “ADOPT: Asynchronous Distributed Constraint Optimization with Quality Guarantees”. In: *Artificial Intelligence* 161.2, pp. 149–180.



PETCU, Adrian and Boi FALTINGS (2005). “A scalable method for multiagent constraint optimization”. In: *IJCAI International Joint Conference on Artificial Intelligence*, pp. 266–271. ISBN: 1045-0823.



VINYALS, Meritxell, Juan A. RODRÍGUEZ-AGUILAR, and Jesus CERQUIDES (2011). “Constructing a unifying theory of dynamic programming DCOP algorithms via the generalized distributive law”. In: *Autonomous Agents and Multi-Agent Systems* 3.22, pp. 439–464. ISSN: 1387-2532. DOI: 10.1007/s10458-010-9132-7.



WEISS, Yair (Jan. 2000). “Correctness of Local Probability Propagation in Graphical Models with Loops”. In: *Neural Comput.* 12.1, pp. 1–41. ISSN: 0899-7667. DOI: 10.1162/089976600300015880. URL: <http://dx.doi.org/10.1162/089976600300015880>.



YOKOO, M. (2001). *Distributed Constraint Satisfaction: Foundations of Cooperation in Multi-Agent Systems*. Springer.

References (cont.)



ZHANG, W., G. WANG, Z. XING, and L. WITTENBURG (2005). “Distributed stochastic search and distributed breakout: properties, comparison and applications to constraint optimization problems in sensor networks.”. In: *Journal of Artificial Intelligence Research (JAIR)* 161.1-2, pp. 55–87.



ZHANG, Weixiong, Guandong WANG, Zhao XING, and Lars WITTENBURG (2003). “A Comparative Study of Distributed Constraint Algorithms”. In: *Distributed Sensor Networks: A Multiagent Perspective*. Ed. by Victor LESSER, Charles L. ORTIZ, and Milind TAMBE. Boston, MA: Springer US, pp. 319–338.