

# Distributed Constraint Processing

## An Introduction

Gauthier Picard

MINES Saint-Étienne  
LaHC UMR CNRS 5516

[gauthier.picard@emse.fr](mailto:gauthier.picard@emse.fr)

— Some contents taken from OPTMAS 2011 and OPTMAS-DCR 2014 Tutorials—

# Contents

## Introduction

- Constraint Satisfaction Problems
- Multi-Agent Approaches to DisCSP

## ABT and Extensions

- Asynchronous Algorithms for DisCSP
- ABT
- AWCS

## Distributed Local Search

- Classical Centralised LS Algorithms
- Distributed Breakout Algorithm (DBA)
- Environment, Reactive rules and Agents (ERA)

## Synthesis

- Panorama
- Using Distributed Problem Solving

# Introduction

## Motivations

- Multi-agent systems are a way to model decentralised problem solving (privacy, distribution)
  - Agents, having personal goals and constraints, negotiate as to reach a global equilibrium
- ⇒ distributed problem solving using agents

## Approaches

- Classical CSP solver extensions
- Classical local search solver extensions

# Cooperative Decentralized Decision Making

## ■ Decentralised Decision Making

- ▶ Agents have to coordinate to perform best actions

## ■ Cooperative settings

- ▶ Agents form a **team** → best actions for the **team**

## ■ Why DDM in cooperative settings is important

- ▶ Surveillance (target tracking, coverage)
- ▶ Robotics (cooperative exploration)
- ▶ Autonomous cars (cooperative traffic management)
- ▶ Scheduling (meeting scheduling)
- ▶ Rescue Operation (task assignment)

# Distributed Constraint Optimisation Problems (DCOPs) for DDM

## Why DCOPs for Cooperative DDM?

### ■ Well defined problem


- ▶ Clear formulation that captures most important aspects
- ▶ Many solution techniques
  - ▶ Optimal: ABT, ADOPT, DPOP, ...
  - ▶ Approximate: DSA, MGM, Max-Sum, ...

### ■ Solution techniques can handle large problems

- ▶ compared for example to sequential decision making (MDP, POMDP)

# Modeling Problems as DCOP


## Target Tracking



### ■ Why decentralize

- Robustness to failure and message loss

## Meeting Scheduling




### ■ Why decentralize


- Privacy

# Target Tracking as a DCOP

- Variables  $\rightarrow$  Cameras
- Domains  $\rightarrow$  Camera actions
  - ▶ look left, look right
- Constraints
  - ▶ Overlapping cameras
  - ▶ Related to targets
    - ▶ Diabolik, Eva
- Maximise sum of constraints

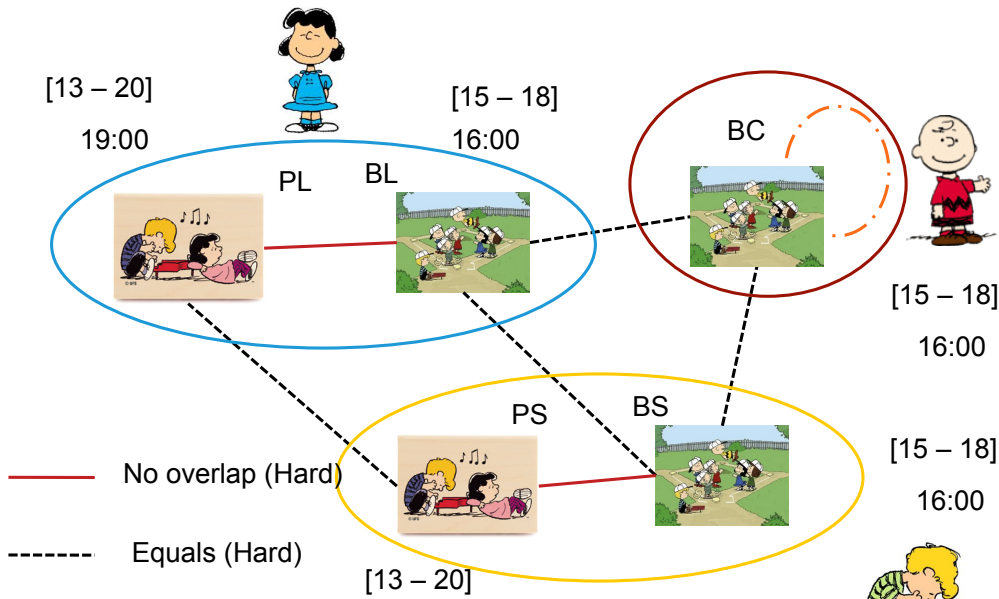


## Meeting Scheduling as a DCOP





# Meeting Scheduling as a DCOP



# Benchmarking problems

## Motivations


- Analysis of complexity and optimality is not enough
- Need to empirically evaluate algorithms on the same problem

## Graph coloring

- Simple to formalise very hard to solve
  - ▶ Well known parameters that influence complexity
    - ▶ Number of nodes, number of colors, density (number of link/number of nodes)
- Many versions of the problem
  - ▶ CSP, MaxCSP, COP

# Graph Coloring


- Network of nodes
- Nodes can take on various colors
- Adjacent nodes should not have the same color
  - ▶ If it happens this is a conflict




**CSP**

# Graph Coloring


- Network of nodes
- Nodes can take on various colors
- Adjacent nodes should not have the same color
  - ▶ If it happens this is a conflict



**CSP**




**Yes**




**No**

## Graph Coloring - MaxCSP


- Optimization Problem
- Natural extension of CSP
- Minimise number of conflicts



0




-1




-4

## Weighted Graph Coloring - COP


- Optimization Problem
- Conflicts have a weight
- Maximise the sum of weights of violated constraints



**COP**



**-2**



**-1**

# Constraint Satisfaction Problems [DECHTER, 2003]

## Definition (CSP)

A CSP is a triplet  $\langle X, D, C \rangle$  such as:

- $X = \{x_1, \dots, x_n\}$  is the set of *variables* to instantiate.
- $D = \{D_1, \dots, D_m\}$  is the set of *domains*. Each variable  $x_i$  is related to a domain of value.
- $C = \{c_1, \dots, c_k\}$  is the set of *constraints*, which are relations between some variables from  $X$  that constrain the values the variables can be simultaneously instantiated to.

## Definition (Solution to a CSP)

A solution to a CSP is a complete assignment of values from  $D$  to variables from  $X$  such that every constraint in  $C$  is satisfied.

# Issues in CSP

## Classical CSPs

- Constraint satisfaction is NP-complete in general
- Constraints are generally expressed as binary constraints
- The topology of a constraint-based problem can be represented by a *constraint network*, in which vertexes represent variables and edges represent binary constraints between variables

## Extensions

- Distribution : variables, constraints
  - ▶ ex.: constraint  $c_i$  belongs to stakeholder  $j$ ,  $\phi(c_i) = j$  (or  $belongs(c_i, j)$ )
- Dynamics : adding removing variables and/or constraints at runtime



# Multi-Agent Approaches to CSP

- **Complete and asynchronous solvers** for combinatorial problems, within the DisCSP framework, such as Asynchronous Backtracking (ABT) or Asynchronous Weak-Commitment Search (AWCS)
- **Distributed local search** methods, such as Distributed Breakout Algorithm (DBA) or Environment, Reactive rules and Agents (ERA) approach

# Asynchronous Algorithms for DisCSP

## Idea

- Inspired by classical centralised algorithms to solve CSP
- Each agent is responsible for assigning one (or several) variables
- Agents propose values to some other agents (depending on the organisation i.e. constraint network)

## Main algorithm: Asynchronous backtracking (ABT) [Yokoo, 2001]

- Agents will perform a distributed version of the backtracking procedure
- ABT is complete
- Extensions exist to handle dynamics

## Definition (DisCSP or DCSP)

A *DisCSP* (or *DCSP*) is a 5-uplet  $\langle A, X, D, C, \phi \rangle$  where  $\langle X, D, C \rangle$  is a CSP,  $A$  is a set of agents and  $\phi : X \mapsto A$  is a function assigning variables from  $X$  to agents from  $A$ .

# Centralised Backtracking

```

 $i \leftarrow 0$ 
 $D'_i \leftarrow D_i$ 
while  $0 \leq i < n$  do
     $x_i \leftarrow \text{null}$ 
     $ok? \leftarrow \text{false}$ 
    while not ok? and  $D'_i$  not empty do
         $a \leftarrow$  a value from  $D'_i$ 
        remove  $a$  from  $D'_i$ 
        if  $a$  is in conflict with  $\{x_0, \dots, x_{i-1}\}$  then
             $x_i \leftarrow a$ 
             $ok? \leftarrow \text{true}$ 
        end
    end
    if  $x_i$  is null then backtrack
         $i \leftarrow i - 1$ 
    else
         $i \leftarrow i + 1$ 
         $D'_i \leftarrow D_i$ 
    end
end


```

# Asynchronous Backtracking (ABT) [Yokoo, 2001]

- First complete asynchronous algorithm for DisCSP solving
- Asynchronous:
  - ▶ All agents active, take a value and inform
  - ▶ No agent has to wait for other agents
- Total order among agents: to avoid cycles
  - ▶  $i < j < k$  means that:  $i$  more priority than  $j$ ,  $j$  more priority than  $k$
- Constraints are directed, following total order
- ABT plays in asynchronous distributed context the same role as backtracking in centralized


# ABT: Directed Constraints

- Directed: from higher to lower priority agents
- Higher priority agent ( $j$ ) informs the lower one ( $k$ ) of its assignment
- Lower priority agent ( $k$ ) evaluates the constraint with its own assignment
  - ▶ If permitted, no action
  - ▶ else it looks for a value consistent with  $j$ 
    - ▶ If it exists,  $k$  takes that value
    - ▶ else, the agent view of  $k$  is a nogood, backtrack



## ABT: Directed Constraints

- Directed: from higher to lower priority agents
- Higher priority agent ( $j$ ) informs the lower one ( $k$ ) of its assignment
- Lower priority agent ( $k$ ) evaluates the constraint with its own assignment
  - ▶ If permitted, no action
  - ▶ else it looks for a value consistent with  $j$ 
    - ▶ If it exists,  $k$  takes that value
    - ▶ else, the agent view of  $k$  is a nogood, backtrack



generates nogoods: eliminate values of  $k$

# ABT: Nogoods

## Definition (Nogood)

Conjunction of (variable, value) pairs of higher priority agents, that removes a value of the current one

## Example

- $x \neq y, d_x = d_y = \{a, b\}, x$  higher than  $y$
- When  $[x \leftarrow a]$  arrives to  $y$ , this agent generates the nogood  $[x = a \Rightarrow y \neq a]$  that removes value  $a$  of  $d_y$
- If  $x$  changes value, when  $[x \leftarrow b]$  arrives to  $y$ , the nogood  $[x = a \Rightarrow y \neq a]$  is eliminated, value  $a$  is again available and a new nogood removing  $b$  is generated

# ABT: Nogood Resolution

- When all values of variable  $y$  are removed, the conjunction of the left-hand sides of its nogoods is also a nogood
- **Resolution:** the process of generating the new nogood

## Example

- $x \neq y, z \neq y, d_x = d_y = d_z = \{a, b\}, x, z$  higher than  $y$

$$x = a \Rightarrow y \neq a$$

$$z = b \Rightarrow y \neq b$$

$x = a \wedge z = b$  is a nogood


$x = a \Rightarrow z \neq b$  (assuming  $x$  higher than  $z$ )



## How ABT works


- **ABT agents:** asynchronous action, spontaneous assignment
- **Assignment:**  $j$  takes value  $a$ ,  $j$  informs lower priority agents
- **Backtrack:**  $k$  has no consistent values with high priority agents,  $k$  resolves nogoods and sends a backtrack message
- **New links:**  $j$  receives a nogood mentioning  $i$ , unconnected with  $j$ ;  $j$  asks  $i$  to set up a link
- **Stop:** “no solution” detected by an agent, stop
- **Solution:** when agents are silent for a while (quiescence), every constraint is satisfied → solution; detected by specialized algorithms

# ABT: Messages




# ABT: Messages

- $Ok?(i \rightarrow k, a)$ :
  - $i$  informs  $k$  that it takes value  $a$




# ABT: Messages

- $Ok?(i \rightarrow k, a)$ :
  - ▶  $i$  informs  $k$  that it takes value  $a$
- $Ngd(k \rightarrow j, i = a \Rightarrow j \neq b)$ 
  - ▶ all  $k$  values are forbidden
  - ▶  $k$  requests  $j$  to backtrack
  - ▶  $k$  forgets  $j$  value
  - ▶  $k$  takes some value
  - ▶  $j$  may detect obsolescence




# ABT: Messages

- $Ok?(i \rightarrow k, a)$ :
  - ▶  $i$  informs  $k$  that it takes value  $a$
- $Ngd(k \rightarrow j, i = a \Rightarrow j \neq b)$ 
  - ▶ all  $k$  values are forbidden
  - ▶  $k$  requests  $j$  to backtrack
  - ▶  $k$  forgets  $j$  value
  - ▶  $k$  takes some value
  - ▶  $j$  may detect obsolescence
- $Addl(j \rightarrow i)$ :
  - ▶ set a link from  $i$  to  $j$ , to know  $i$  value



# ABT: Messages

- $Ok?(i \rightarrow k, a)$ :
  - ▶  $i$  informs  $k$  that it takes value  $a$
- $Ngd(k \rightarrow j, i = a \Rightarrow j \neq b)$ 
  - ▶ all  $k$  values are forbidden
  - ▶  $k$  requests  $j$  to backtrack
  - ▶  $k$  forgets  $j$  value
  - ▶  $k$  takes some value
  - ▶  $j$  may detect obsolescence
- $Addl(j \rightarrow i)$ :
  - ▶ set a link from  $i$  to  $j$ , to know  $i$  value
- Stop:
  - ▶ there is no solution



# ABT Procedures

**when received** (*ok?*,  $(x_j, d_j)$ ) **do** — (i)

  revise *agent\_view*;

**check\_agent\_view**;

**end do**;

**when received** (*nogood*,  $x_j$ , *nogood*) **do** — (ii)

  record *nogood* as a new constraint;

**when** *nogood* contains an agent  $x_k$  that is not its neighbor

**do** request  $x_k$  to add  $x_i$  as a neighbor,

    and add  $x_k$  to its neighbors; **end do**;

*old\_value*  $\leftarrow$  *current\_value*; **check\_agent\_view**;

**when** *old\_value* = *current\_value* **do**

    send (*ok?*,  $(x_j, \text{current\_value})$ ) to  $x_j$ ; **end do**; **end do**;

**procedure** **check\_agent\_view**

**when** *agent\_view* and *current\_value* are not consistent **do**

**if** no value in  $D_i$  is consistent with *agent\_view* **then** **backtrack**;

**else** select  $d \in D_i$  where *agent\_view* and  $d$  are consistent;

*current\_value*  $\leftarrow$   $d$ ;

    send (*ok?*,  $(x_i, d)$ ) to neighbors; **end if**; **end do**;

**procedure** **backtrack**

  generate a nogood  $V$  — (iii)

**when**  $V$  is an empty nogood **do**

    broadcast to other agents that there is no solution,

    terminate this algorithm; **end do**;




  select  $(x_j, d_j)$  where  $x_j$  has the lowest priority in a nogood;

  send (*nogood*,  $x_i$ ,  $V$ ) to  $x_j$ ;

  remove  $(x_j, d_j)$  from *agent\_view*;

**check\_agent\_view**;

**Algorithm 2:** ABT Procedures



# ABT: Correctness and Completeness

## ■ Correctness

- ▶ silent network  $\Leftrightarrow$  all constraints are satisfied

## ■ Completeness

- ▶ ABT performs an exhaustive traversal of the search space
- ▶ Parts not searched: those eliminated by nogoods
- ▶ Nogoods are legal: logical consequences of constraints
- ▶ Therefore, either there is no solution  $\Rightarrow$  ABT generates the empty nogood, or it finds a solution if exists



## ABT: Remarks

- Fixed ordered organisation
  - ▶ Agents only communicate with agents with lower priority for *ok?*
  - ▶ Agents only communicate with the agent with direct higher priority for *nogood*
- No termination procedure is given (but it is easily implemented using Dijkstra's tokens)
- Really distributable
- What if  $x_0$  disappears?...

## Extensions and Filiation

- **Changing ordering** in every conflict with AWCS [YOKOO, 2001]
- Satisfaction → **Optimisation** with ADOPT (Asynchronous B&B) [MODI et al., 2005] or APO [MAILLER and LESSER, 2006]
- **Adding new agents** at runtime in DynAPO [MAILLER, 2005]

# Asynchronous Weak-Commitment Search (AWCS) [Yokoo, 2001]

procedure **check\_agent\_view**

```





when agent.view and current.value are not consistent do
  if no value in  $D_i$  is consistent with agent.view then backtrack;
  else select  $d \in D_i$  where agent.view and  $d$  are consistent
    and  $d$  minimizes the number of constraint violations
    with lower priority agents; — (i)
    current.value  $\leftarrow d$ ;
    send (ok?, ( $x_i$ ,  $d$ , current.priority)) to neighbors;
  end if; end do;
  
```

procedure **backtrack**

```

generate a nogood  $V$ ;
when  $V$  is an empty nogood do
  broadcast to other agents that there is no solution,
  terminate this algorithm; end do;
when  $V$  is a new nogood do — (ii)
  send  $V$  to the agents in the nogood;
  current.priority  $\leftarrow 1 + p_{max}$ ,
  where  $p_{max}$  is the maximal priority value of neighbors;
  select  $d \in D_i$  where agent.view and  $d$  are consistent,
  and  $d$  minimizes the number of constraint violations
  with lower priority agents;
  current.value  $\leftarrow d$ ;
  send (ok?, ( $x_i$ ,  $d$ , current.priority)) to neighbors; end do;
  
```

**Algorithm 3:** AWCS Procedures



# Distributed Local Search Approaches

## Local Search (LS)

- LS algorithms explore the search space from state to state
- Always tend to improve the current state of the system
- Can naturally handle dynamics (adding constraints, changing values)
- Time efficient
- Not complete and require some subtle parameter tuning

choose an initial assignment  $s(0)$

**while**  $s(t)$  *not terminal* **do**

    | select an acceptable move  $m(t)$  to another assignment

    | apply move  $m(t)$  to reach  $s(t + 1)$

    |  $t := t + 1$

**end**

**Algorithm 4:** A generic centralised local search algorithm

# Classical Centralised LS Algorithms

## Common points

- Initial point (ex: randomly chosen)
- Termination criterion (ex: limit time,  $\delta$  improvement)
- Acceptable move (ex:  $+\epsilon$ )

## Famous LS Methods

- **Tabu search** [GLOVER and LAGUNA, 1997]
- **Simulated annealing** [KIRKPATRICK et al., 1983]
- **Iterative Breakout method** [MORRIS, 1993]

# Distributed Breakout Algorithm (DBA)

```

wait_ok? mode — (i)
when received (ok?,  $x_j$ ,  $d_j$ ) do
  add ( $x_j$ ,  $d_j$ ) to agent_view;
  when received ok? messages from all neighbors do
    send_improve;
    goto wait_improve mode; end do;
  goto wait_ok mode; end do;

procedure send_improve
  current_eval  $\leftarrow$  evaluation value of current_value;
  my_improve  $\leftarrow$  possible maximal improvement;
  new_value  $\leftarrow$  the value which gives the maximal improvement;
  send (improve,  $x_i$ , my_improve, current_eval) to neighbors;

```

```

wait_improve? mode — (ii)
when received (improve,  $x_j$ , improve, eval) do
  record this message;
  when received improve? messages from all neighbors do
    send_ok; clear agent_view;
    goto wait_ok mode; end do;
  goto wait_improve mode; end do;





```

```

procedure send_ok
  when its improvement is largest among neighbors do
    current_value  $\leftarrow$  new_value; end do;
  when it is in a quasi-local-minimum do
    increase the weights of constraint violations; end do;
  send (ok?,  $x_i$ , current_value) to neighbors;

```

**Algorithm 5:** DBA Message Handler



## Distributed Breakout Algorithm (DBA) (cont.)

### Principles of DBA [Yokoo, 2001]

#### ■ Distribution difficulties:

- (i) if two neighbouring agents concurrently change their value, the *system may oscillate*
- (ii) detecting the fact that the whole system is trapped in local minimum requires the agents to *globally exchange data*

#### ■ DBA answers:

- (i) for a given neighbourhood, only the agent that can *maximally improve* the evaluation value is given the right to change its value
- (ii) agents only detects *quasi-local-minimum*, which is a weaker local-minimum that can be detected only by local interactions

## Distributed Breakout Algorithm (DBA) (cont.)

### Remarks

- Distributed version of the iterative breakout algorithm
- Two-mode behaviour alternating between exchange of potential improvement and exchange of assignments
- ✓ There is no order over the agents society → neighbourhoods
- The system halts if a solution is found or if the weight of constraints have reached a *predefined upper bound*
  - the **only** difficult parameter to set
- ✗ DBA is not complete
- ✓ DBA is able to detect the termination or a global solution only by reasoning on local data.

# Environment, Reactive rules and Agents (ERA) [Liu et al., 2002]

## Components

- A discrete grid **environment**, that is used as a communication medium
- **Agents** that evolves in some regions of the grid (their domain)
  - ▶ Agents move *synchronously*
  - ▶ Agents cannot move in the domain of other agents, but can mark it with the number of potential conflicts
  - ▶ These marks represents therefore the number of violated constraints if an agent chooses the marked cell
- **Rules** (*moves*) that agent follow to reach an equilibrium
  - ▶ 3 possible actions
    - ▶ *least-move*: the next cell is the one with minimum cost
    - ▶ *better-move*: the next cell is randomly chosen and if it has less conflicts than the actual one the agent moves else the agent rests
    - ▶ *random-move*: the next cell is randomly chosen
  - ▶ A decision consists in a random Monte-Carlo choice of the action to perform

$t \leftarrow 0$

initialise the grid to 0 violation in each cell; **foreach** agent  $i$  **do**

    | randomly move to a cell of row  $i$

**end**

**while**  $t < t_{max}$  and no solution **do**

    | **foreach** agent  $i$  **do**

        | select a move behaviour

        | compute new position

        | decrease markers in all cells with past violations

    | .. .. .. .. ..

$X_1$	1	2	1	1	0	0
$X_2$	0	1	0	0		
$X_3$	0	1	1	1	1	

$a_1$  *least-move*

$X_1$	1	2	1	1	0	0
$X_2$	0	0	0	0		
$X_3$	0	0	0	0	1	

$X_1$	1	2	1	1	0	0
$X_2$	0	1	0	0		
$X_3$	0	1	1	1	1	

$a_1$  *better-move*

$X_1$	1	1	1	1	0	0
$X_2$	0	0	1	0		
$X_3$	0	0	1	1	1	



# Environment, Reactive rules and Agents (ERA) [Liu et al., 2002] (cont.)

## Remarks

- The environment is the communication medium
  - ✓ There is no asynchronous mechanisms and message handling
  - ✗ Synchronisation point: high synchronous solving process with no benefit from distribution, in case of high connected constraint networks
- ✓ ERA quickly finds assignments close to the solution → repairing issues
- ✗ Redundant usage of random choices: non-guided method, close to random walk, and non complete
- ✗ Termination: ERA requires a time limit ( $t_{max}$ ) (problem-dependant)

# Panorama

Algorithm	Type	Memory	Messages	Remarks
ABT	CSP	Exponential	–	Complete, Static ordering
AWCS	CSP	Exponential	–	Complete (only with exponential space), Reordering, fast
DBA	Max-CSP	Linear	Bounded	Incomplete, Fast
ERA	Max-CSP	Polynomial	n/a	Incomplete, randomness

Table: DCSP and DCOP algorithms

# Using Distributed Problem Solving

## Problem and Environment Characteristics

- Geographic distribution
  - ▶ ex: agents are physically distributed, and solving the whole problem is not possible in a centralised manner
- Constraint network topology
  - ▶ ex: bounded vertex degrees or large constraint graph diameter
- Knowledge encapsulation
  - ▶ ex: *privacy* preserving, limited knowledge
- Dynamics
  - ▶ ex: rather than solving the whole problem again, only repair sub-problems

## Some Applications

- Frequency assignment
- Scheduling
- Resource allocation, Manufacturing control
- Supply chain
- ...

# References



DECHTER, R. (2003). *Constraint Processing*. Morgan Kaufmann.



GLOVER, F. and M. LAGUNA (1997). *Tabu Search*. Kluwer.



KIRKPATRICK, S., C. GELLAT, and M. VECCHI (1983). “Optimization by Simulated Annealing”. In: *Science* 220.4598, pp. 671–680.



LIU, J., H. JING, and Y. Y. TANG (2002). “Multi-agent Oriented Constraint Satisfaction”. In: *Artificial Intelligence* 136.1, pp. 101–144.



MAILLER, R. (2005). “Comparing two approaches to dynamic, distributed constraint satisfaction”. In: *Proceedings of the Fourth International Joint Conference on Autonomous Agents and Multiagent Systems (AAMAS’05)*. ACM Press, pp. 1049–1056.



MAILLER, R. and V. R. LESSER (2006). “Asynchronous Partial Overlay: A New Algorithm for Solving Distributed Constraint Satisfaction Problems”. In: *Journal of Artificial Intelligence Research* 25, pp. 529–576.



MODI, P. J., W. SHEN, M. TAMBE, and M. YOKOO (2005). “ADOPT: Asynchronous Distributed Constraint Optimization with Quality Guarantees”. In: *Artificial Intelligence* 161.2, pp. 149–180.



MORRIS, P. (1993). “The Breakout Method for Escaping from Local Minima”. In: *AAAI*, pp. 40–45.



YOKOO, M. (2001). *Distributed Constraint Satisfaction: Foundations of Cooperation in Multi-Agent Systems*. Springer.